

Ximing Lu

Ph.D. Candidate, University of Washington

✉ ximinglux32@gmail.com | 🏠 gloriaximinglu.github.io | 🐦 @GXiming

EDUCATION

University of Washington

Ph.D. in Computer Science (Advisor: Yejin Choi)

Seattle, WA
2023 – Present

University of Washington

B.S. in Computer Science (Summa Cum Laude)

Seattle, WA
2017 – 2023

EXPERIENCE

NVIDIA Research

Research Scientist, Language and Cognition Research

Santa Clara, CA
2024 – 2026

Allen Institute for AI (AI2)

Pre-Doctoral Young Investigator, Mosaic Team

Seattle, WA
2021 – 2024

Allen Institute for AI (AI2)

Research Intern, Mosaic Team

Seattle, WA
2020 – 2021

University of Washington

Research Assistant, Xlab

Seattle, WA
2019

SELECTED AWARDS & RECOGNITION

- **Best Paper Award**, NAACL 2022 · **Outstanding Paper Award**, EMNLP 2023
- **6,900+ Google Scholar citations** across **85+ publications**; **20+ papers with 100+ citations**; **h-index 34**
- **10+ Oral/Spotlight presentations** at NeurIPS, ICML, ICLR, CVPR, NAACL (incl. top 1% selections)
- Featured in **Science Magazine** (*News from Science*) for *Creativity Index* paper
- *Faith and Fate* paper: **484K views**, 1.3K likes, 385 reposts on Twitter; endorsed by **Yann LeCun** (Turing Award) and **Kevin Murphy** (Google DeepMind)
- *ProRL* paper: **177K views** on Twitter; endorsed by **Nathan Lambert** (AI2); open-sourced in NVIDIA NeMo Gym

OPEN ARTIFACTS

Nemotron 3 Nano Omni: Efficient and Open Multimodal Intelligence · *Cumulative Downloads: 1.8M*

Nemotron-PrismMath: Gradient-based Data Diversification Boosts Generalization in LLM Reasoning · *Downloads: 50K*

SODA Dataset: Million-scale Dialogue with Social Commonsense · *Downloads: 31K+* · *EMNLP 2023 Outstanding Paper*

ProsocialDialog Dataset: A Prosocial Backbone for Conversational Agents · *Downloads: 29K+*

PUBLICATIONS

* = equal contribution. Underlined = X Lu.

2026

Y Lin, S Tang, B Lyu, ..., X Lu, D Acuna, K Yang, H Lin, Y Choi, D Chen, S Arora, C Jin. **Goedel-Prover-V2: Scaling Formal Theorem Proving with Scaffolded Data Synthesis and Self-Correction**. *In Proceedings of the International Conference on Learning Representations*, 2026.

F Wu*, W Xuan*, H Qi*, A Tu, X Lu, LE Li, Y Choi. **DeepSearch: Overcome the Bottleneck of Reinforcement Learning with Verifiable Rewards via Tree-based Search**. *In Proceedings of the International Conference on Learning Representations*, 2026.

Z Wang, J Jung, X Lu, S Diao, E Evans, J Zeng, P Molchanov, Y Choi, J Kautz, Y Dong. **ProfBench: Multi-Domain Rubrics requiring Professional Knowledge to Answer and Judge**. *In Proceedings of the International Conference on Learning Representations*, 2026.

X Lu, D Acuna, J Jung, J Hu, D Zhang, S Diao, ..., J Kautz, Y Dong, Y Choi. **Golden Goose: A Simple Trick to Synthesize Unlimited RLVR Tasks from Unverifiable Internet Text**. *In Proceedings of the International Conference on Machine Learning*, 2026.

D Acuna*, C-H H Yang*, Y Deng*, J Jung*, X Lu*, P Ammanabrolu, H Kim, Y-H Liao, Y Choi. **Long Grounded Thoughts: Synthesizing Visual Problems and Reasoning Chains at Scale**. *In Proceedings of the International Conference on Machine Learning*, 2026.

J Hu, M Liu, [X Lu](#), F Wu, Z Harchaoui, S Diao, Y Choi, P Molchanov, J Yang, J Kautz, Y Dong. **BroRL: Scaling Reinforcement Learning via Broadened Exploration**. *In Proceedings of the International Conference on Machine Learning*, 2026.

H Su*, S Diao*, [X Lu](#), M Liu, J Xu, X Dong, ..., Y Choi, J Kautz, P Molchanov. **ToolOrchestra: Elevating Intelligence via Efficient Model and Tool Orchestration**. *In Proceedings of the International Conference on Machine Learning*, 2026.

S-Y Liu, X Dong, [X Lu](#), S Diao, P Belcak, M Liu, ..., Y Choi, J Kautz, P Molchanov. **GDPO: Group reward-Decoupled Normalization Policy Optimization for Multi-reward RL Optimization**. *In Proceedings of the International Conference on Machine Learning*, 2026.

NVIDIA (incl. [X Lu](#)). **Nemotron 3 Nano Omni: Efficient and Open Multimodal Intelligence**. *NVIDIA Technical Report*, 2026. [link]

H Zhang, M Liu, S Zhang, S Han, J Hu, Z Jin, Y Zhang, S Diao, [X Lu](#), B Xu, Z Yu, J Kautz, Y Dong. **ProRL Agent: Rollout-as-a-Service for RL Training of Multi-Turn LLM Agents**. *NVIDIA Technical Report*, 2026. [link]

2025

M Liu, S Diao, [X Lu](#), J Hu, X Dong, Y Choi, J Kautz, Y Dong. **ProRL: Prolonged Reinforcement Learning Expands Reasoning Boundaries in Large Language Models**. *In Proceedings of the Conference on Neural Information Processing Systems*, 2025.
Featured: 177K views on Twitter/X · Endorsed by Nathan Lambert (AI2)

J Jung, S Han*, [X Lu](#)*, S Hallinan*, D Acuna, S Prabhumoye, M Patwary, M Shoeybi, B Catanzaro, Y Choi. **Prismatic Synthesis: Gradient-based Data Diversification Boosts Generalization in LLM Reasoning**. *In Proceedings of the Conference on Neural Information Processing Systems*, 2025 (**Spotlight**).

[X Lu](#), M Sclar, S Hallinan, N Mireshghallah, J Liu, S Han, A Ettinger, L Jiang, K Chandu, N Dziri, Y Choi. **AI as Humanity's Salieri: Quantifying Linguistic Creativity of Language Models via Systematic Attribution of Machine Text against Web Text**. *In Proceedings of the International Conference on Learning Representations*, 2025 (**Oral, top 1.8%**).
Featured: Science Magazine (News from Science) · CGTN America · Alwire · Ploutos AI

K Chandu*, L Li*, A Awadalla, [X Lu](#), JS Park, J Hessel, L Wang, Y Choi. **CertainlyUncertain: A Benchmark and Metric for Multimodal Epistemic and Aleatoric Awareness**. *In Proceedings of the International Conference on Learning Representations*, 2025.

X Zhou, H Kim*, F Brahman*, L Jiang, H Zhu, [X Lu](#), F Xu, BY Lin, Y Choi, N Mireshghallah, R Le Bras, M Sap. **HAICOSYSTEM: An Ecosystem for Sandboxing Safety Risks in Human-AI Interactions**. *In Proceedings of the Conference on Language Modeling*, 2025.

S Hallinan, J Jung, M Sclar, [X Lu](#), A Ravichander, S Ramnath, Y Choi, SP Karimireddy, N Mireshghallah, X Ren. **The Surprising Effectiveness of Membership Inference with Simple N-Gram Coverage**. *In Proceedings of the Conference on Language Modeling*, 2025.

W Seo*, S Han*, J Jung, B Newman, S Lim, S Lee, [X Lu](#), Y Choi, Y Yu. **Verifying the Verifiers: Unveiling Pitfalls and Potentials in Fact Verifiers**. *In Proceedings of the Conference on Language Modeling*, 2025.

D Acuna, [X Lu](#), J Jung, H Kim, A Kar, S Fidler, Y Choi. **Socratic-MCTS: Test-Time Visual Reasoning by Asking the Right Questions**. *In Proceedings of the Conference on Empirical Methods in Natural Language Processing*, 2025.

JS Park, Z Ma, L Li, C Zheng, C-Y Hsieh, [X Lu](#), K Chandu, Q Kong, N Kobori, A Farhadi, Y Choi, R Krishna. **Synthetic Visual Genome**. *In Proceedings of the Conference on Computer Vision and Pattern Recognition*, 2025.

A Ravichander, J Fisher, T Sorensen, [X Lu](#), Y Lin, M Antoniak, N Mireshghallah, C Bhagavatula, Y Choi. **Information-Guided Identification of Training Data Imprint in (Proprietary) Large Language Models**. *In Proceedings of the North American Chapter of the Association for Computational Linguistics*, 2025.

A Baheti, D Chakraborty, F Brahman, R Le Bras, [X Lu](#), N Dziri, Y Choi, M Riedl, M Sap. **Multi-Attribute Constraint Satisfaction via Language Model Rewriting**. *Transactions on Machine Learning Research*, 2025.

S-Y Liu, X Dong, [X Lu](#), S Diao, M Liu, M-H Chen, H Yin, Y-C F Wang, K-T Cheng, Y Choi, J Kautz, P Molchanov. **DLER: Doing Length Penalty Right — Incentivizing More Intelligence per Token via Reinforcement Learning**. *NVIDIA Technical Report*, 2025. [link]

NVIDIA (incl. [X Lu](#)). **NVIDIA Nemotron Nano 2: An Accurate and Efficient Hybrid Mamba-Transformer Reasoning Model**. *NVIDIA Technical Report*, 2025. [link]

M Liu, S Diao, J Hu, [X Lu](#), X Dong, ..., Y Choi, J Kautz, Y Dong. **Scaling Up RL: Unlocking Diverse Reasoning in LLMs via Prolonged Training**. *NVIDIA Technical Report*, 2025. [link]

2024

L Jiang, K Rao*, S Han*, A Ettinger, F Brahman, S Kumar, N Mireshghallah, [X Lu](#), M Sap, Y Choi, N Dziri. **WildTeaming at Scale: From In-the-Wild Jailbreaks to (Adversarially) Safer Language Models**. *In Proceedings of the Conference on Neural Information Processing Systems*, 2024.

T Sorensen, J Moore, J Fisher, M Gordon, N Mireshghallah, CM Rytting, A Ye, L Jiang, [X Lu](#), N Dziri, T Althoff, Y Choi. **Position: A Roadmap to Pluralistic Alignment**. *In Proceedings of the International Conference on Machine Learning*, 2024.

L Qiu, L Jiang, [X Lu](#), M Sclar, V Pyatkin, C Bhagavatula, B Wang, Y Kim, Y Choi, N Dziri, X Ren. **Phenomenal Yet Puzzling: Testing Inductive Reasoning Capabilities of Language Models with Hypothesis Refinement**. *In Proceedings of the International Conference on Learning Representations*, 2024 (**Oral, top 1.2%**).

BY Lin, A Ravichander, [X Lu](#), N Dziri, M Sclar, K Chandu, C Bhagavatula, Y Choi. **The Unlocking Spell on Base LLMs: Rethinking Alignment via In-Context Learning**. *In Proceedings of the International Conference on Learning Representations*, 2024.

- P West*, X Lu*, N Dziri*, F Brahman*, L Li*, JD Hwang, L Jiang, J Fisher, A Ravichander, K Chandu, B Newman, PW Koh, A Ettinger, Y Choi. **The Generative AI Paradox: “What It Can Create, It May Not Understand.”** In *Proceedings of the International Conference on Learning Representations*, 2024.
- A Baheti, X Lu, F Brahman, R Le Bras, M Sap, M Riedl. **Leftover Lunch: Advantage-based Offline Reinforcement Learning for Language Models.** In *Proceedings of the International Conference on Learning Representations*, 2024.
- S Ramnath, B Joshi, S Hallinan, X Lu, LH Li, A Chan, J Hessel, Y Choi, X Ren. **Tailoring Self-Rationalizers with Multi-Reward Distillation.** In *Proceedings of the International Conference on Learning Representations*, 2024.
- T Sorensen, L Jiang, JD Hwang, S Levine, V Pyatkin, P West, N Dziri, X Lu, K Rao, C Bhagavatula, M Sap, J Tasioulas, Y Choi. **Value Kaleidoscope: Engaging AI with Pluralistic Human Values, Rights, and Duties.** In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024 (**Oral, top 3%**).
- J Jung, P West, L Jiang, F Brahman, X Lu, J Fisher, T Sorensen, Y Choi. **Impossible Distillation for Paraphrasing and Summarization: How to Make High-quality Lemonade out of Small, Low-quality Models.** In *Proceedings of the North American Chapter of the Association for Computational Linguistics*, 2024.
- J Fisher, X Lu, J Jung, L Jiang, Z Harchaoui, Y Choi. **JAMDEC: Unsupervised Authorship Obfuscation using Constrained Decoding over Small Language Models.** In *Proceedings of the North American Chapter of the Association for Computational Linguistics*, 2024.
- J Jung, X Lu, L Jiang, F Brahman, P West, PW Koh, Y Choi. **Information-Theoretic Distillation for Reference-less Summarization.** In *Proceedings of the Conference on Language Modeling*, 2024.
- J Fisher*, S Hallinan*, X Lu, M Gordon, Z Harchaoui, Y Choi. **StyleRemix: Interpretable Authorship Obfuscation via Distillation and Perturbation of Style Elements.** In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, 2024.
- H Li, Y Ning, Z Liao, S Wang, XL Li, X Lu, W Zhao, F Brahman, Y Choi, X Ren. **In Search of the Long-Tail: Systematic Generation of Long-Tail Inferential Knowledge via Logical Rule Guided Search.** In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, 2024.
- J Lee, X Lu, J Hessel, F Brahman, Y Yu, Y Bisk, Y Choi, S Gabriel. **How to Train Your Fact Verifier: Knowledge Transfer with Multimodal Open Models.** In *Findings of the Conference on Empirical Methods in Natural Language Processing*, 2024.

2023

- N Dziri*, X Lu*, M Sclar*, XL Li, L Jiang, BY Lin, P West, C Bhagavatula, R Le Bras, JD Hwang, S Sanyal, S Welleck, X Ren, A Ettinger, Z Harchaoui, Y Choi. **Faith and Fate: Limits of Transformers on Compositionality.** In *Proceedings of the Conference on Neural Information Processing Systems*, 2023 (**Spotlight**).
Featured: 484K views on Twitter/X · Reposted by Yann LeCun · Reposted by Kevin Murphy · Endorsed by Vitaly Kurin (NVIDIA) · YouTube (Automata Learning Lab) · YouTube (mardin mardin)
- JS Park, J Hessel, K Chandu, PP Liang, X Lu, P West, Y Yu, Q Huang, J Gao, A Farhadi, Y Choi. **Localized Symbolic Knowledge Distillation for Visual Commonsense Models.** In *Proceedings of the Conference on Neural Information Processing Systems*, 2023.
- S Welleck*, X Lu*, P West, F Brahman, T Shen, D Khashabi, Y Choi. **Generating Sequences by Learning to Self-Correct.** In *Proceedings of the International Conference on Learning Representations*, 2023.
- H Kim, J Hessel, L Jiang, P West, X Lu, Y Yu, P Zhou, R Le Bras, M Alikhani, G Kim, M Sap, Y Choi. **SODA: Million-scale Dialogue Distillation with Social Commonsense Contextualization.** In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, 2023 (**Outstanding Paper Award**).
- X Lu, F Brahman, P West, J Jung, K Chandu, A Ravichander, L Qin, P Ammanabrolu, L Jiang, S Ramnath, N Dziri, J Fisher, BY Lin, S Hallinan, X Ren, S Welleck, Y Choi. **Inference-Time Policy Adapters (IPA): Tailoring Extreme-Scale LMs without Fine-tuning.** In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, 2023.
- P West, R Le Bras, T Sorensen, BY Lin, L Jiang, X Lu, K Chandu, J Hessel, A Baheti, C Bhagavatula, Y Choi. **NovaCOMET: Open Commonsense Foundation Models with Symbolic Knowledge Distillation.** In *Findings of the Conference on Empirical Methods in Natural Language Processing*, 2023.
- S Hallinan, F Brahman, X Lu, J Jung, S Welleck, Y Choi. **STEER: Unified Style Transfer with Expert Reinforcement.** In *Findings of the Conference on Empirical Methods in Natural Language Processing*, 2023.
- V Pyatkin, JD Hwang, V Srikumar, X Lu, L Jiang, Y Choi, C Bhagavatula. **ClarifyDelphi: Reinforced Clarification Questions with Defeasibility Rewards for Social and Moral Situations.** In *Proceedings of the Annual Meeting of the Association for Computational Linguistics*, 2023.
- C Bhagavatula*, JD Hwang*, D Downey*, R Le Bras*, X Lu*, L Qin, K Sakaguchi, S Swayamdipta, P West*, Y Choi*. **I2D2: Inductive Knowledge Distillation with NeuroLogic and Self-Imitation.** In *Proceedings of the Annual Meeting of the Association for Computational Linguistics*, 2023.
- Y Yu*, J Chung*, H Yun, J Hessel, JS Park, X Lu, R Zellers, P Ammanabrolu, R Le Bras, G Kim, Y Choi. **Fusing Pre-trained Language Models with Multimodal Prompts through Reinforcement Learning.** In *Proceedings of the Conference on Computer Vision and Pattern Recognition*, 2023.

2022

- X Lu, S Welleck, P West, L Jiang, J Kasai, D Khashabi, R Le Bras, L Qin, Y Yu, R Zellers, NA Smith, Y Choi. **NeuroLogic A^{*}esque Decoding: Constrained Text Generation with Lookahead Heuristics**. In *Proceedings of the North American Chapter of the Association for Computational Linguistics*, 2022 (**Best Paper Award**).
- P West, C Bhagavatula, J Hessel, JD Hwang, L Jiang, R Le Bras, X Lu*, S Welleck*, Y Choi. **Symbolic Knowledge Distillation: from General Language Models to Commonsense Models**. In *Proceedings of the North American Chapter of the Association for Computational Linguistics*, 2022 (**Oral**).
- Y Zhao, J Hessel, Y Yu, X Lu, R Zellers, Y Choi. **Connecting the Dots between Audio and Text without Parallel Data through Visual Knowledge Transfer**. In *Proceedings of the North American Chapter of the Association for Computational Linguistics*, 2022.
- X Lu, S Welleck, J Hessel, L Jiang, L Qin, P West, P Ammanabrolu, Y Choi. **Quark: Controllable Text Generation with Reinforced Unlearning**. In *Proceedings of the Conference on Neural Information Processing Systems*, 2022 (**Oral**).
- S Welleck*, J Liu*, X Lu, H Hajishirzi, Y Choi. **NaturalProver: Grounded Mathematical Proof Generation with Language Models**. In *Proceedings of the Conference on Neural Information Processing Systems*, 2022.
- R Zellers, J Lu, X Lu, Y Yu, Y Zhao, M Salehi, A Kusupati, J Hessel, A Farhadi, Y Choi. **MERLOT Reserve: Multimodal Neural Script Knowledge through Vision and Language and Sound**. In *Proceedings of the Conference on Computer Vision and Pattern Recognition*, 2022 (**Oral**).
- J Liu, A Liu, X Lu, S Welleck, P West, R Le Bras, Y Choi, H Hajishirzi. **Generated Knowledge Prompting for Commonsense Reasoning**. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics*, 2022.
- H Kim, Y Yu, L Jiang, X Lu, D Khashabi, G Kim, Y Choi, M Sap. **ProsocialDialog: A Prosocial Backbone for Conversational Agents**. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, 2022.
- J Liu, S Hallinan, X Lu, P He, S Welleck, H Hajishirzi, Y Choi. **RAINIER: Reinforced Knowledge Introspector for Commonsense Question Answering**. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, 2022.
- J Kasai, K Sakaguchi, R Le Bras, H Peng, X Lu, D Radev, Y Choi, NA Smith. **Twist Decoding: Diverse Generators Guide Each Other**. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, 2022.
- S Mehta, X Lu, W Wu, D Weaver, H Hajishirzi, JG Elmore, LG Shapiro. **End-to-End Diagnosis of Breast Biopsy Images with Transformers**. *Medical Image Analysis*, 2022.

2021

- R Zellers*, X Lu*, J Hessel*, Y Yu, JS Park, J Cao, A Farhadi, Y Choi. **MERLOT: Multimodal Neural Script Knowledge Models**. In *Proceedings of the Conference on Neural Information Processing Systems*, 2021 (**Oral, top 1%**).
- A Liu, M Sap, X Lu, S Swayamdipta, C Bhagavatula, NA Smith, Y Choi. **DExperts: On-the-Fly Controlled Text Generation with Experts and Anti-Experts**. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics*, 2021 (**Oral**).
- P West, X Lu, A Holtzman, C Bhagavatula, J Hwang, Y Choi. **Reflective Decoding: Beyond Unidirectional Generation with Off-the-Shelf Language Models**. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics*, 2021.
- Y Dong, C Bhagavatula, X Lu, JD Hwang, A Bosselut, JCK Cheung, Y Choi. **On-the-Fly Attention Modulation for Neural Generation**. In *Findings of the Annual Meeting of the Association for Computational Linguistics*, 2021.
- X Lu, P West, R Zellers, R Le Bras, C Bhagavatula, Y Choi. **NeuroLogic Decoding: (Un)supervised Neural Text Generation with Predicate Logic Constraints**. In *Proceedings of the North American Chapter of the Association for Computational Linguistics*, 2021.
- J Da, R Le Bras, X Lu, Y Choi, A Bosselut. **Analyzing Commonsense Emergence in Few-shot Knowledge Models**. In *Proceedings of Automated Knowledge Base Construction*, 2021.
- X Lu, S Mehta, TT Brunyé, DL Weaver, JG Elmore, LG Shapiro. **Analysis of Regions of Interest and Distractor Regions in Breast Biopsy Images**. *IEEE EMBS International Conference on Biomedical and Health Informatics*, 2021.
- S Mehta, N Nuechterlein, E Mercan, B Li, S Nofallah, W Wu, X Lu, A Caspi, M Rastegari, J Elmore, H Hajishirzi, L Shapiro. **Applications of the ESPNet Architecture in Medical Imaging**. *Book chapter, State of the Art in Neural Networks and their Applications*, Elsevier (pp. 117–131), 2021.

Pre-prints (Under Review)

- B Cui, X Lu, J Jung, SN Akter, H Kim, Y Qu, D Acuna, S Prabhumoye, Y Choi, P Ammanabrolu. **Introspective X Training: Feedback Conditioning Improves Scaling Across all LLM Training Stages**. *arXiv:2605.20285*, 2026. [link]
- J Jung, H Kim, B Cui, X Lu, D Acuna, P Ammanabrolu, Y Choi. **DeltaPrompts: Escaping the Zero-Delta Trap in Multimodal Distillation**. *arXiv:2605.15532*, 2026. [link]
- A Hatamizadeh, S Prabhumoye, I Gitman, X Lu, S Han, W Ping, Y Choi, J Kautz. **iGRPO: Self-Feedback-Driven LLM Reasoning**. *arXiv:2602.09000*, 2026. [link]
- F Wu, W Xuan, X Lu, M Liu, Y Dong, Z Harchaoui, Y Choi. **The Invisible Leash: Why RLVR May or May Not Escape Its Origin**. *arXiv:2507.14843*, 2025. [link]
- ZJ Hou, BA Zhang, Y Lu, BK Baghel, A Brei, X Lu, M Jiang, F Brahman, S Chaturvedi, H-S Chang, D Khashabi, XL Li. **CreativityPrism: A Holistic Evaluation Framework for Large Language Model Creativity**. *arXiv:2510.20091*, 2025. [link]

MEDIA COVERAGE

AI as Humanity's Salieri (ICLR 2025)

- **Science (AAAS), News from Science** — “AI writing is improving, but it still can’t match human creativity,” by Matthew Hutson. [link]
- **CGTN America** — *World Today* segment “Can AI Match Human Creativity?”, by Mark Niu. [link]
- **Alwire** — “Despite Rapid Advances, Studies Show AI Can’t Match Human Creativity,” by Ali Azhar. [link]
- **Ploutos AI** — Featured research stream (AI talks, demos, and community discussion). [link]

Faith and Fate (NeurIPS 2023)

- **Twitter/X** — post about the paper reached **484,000 views**. [link]
- **Yann LeCun** (Turing Award; former Chief AI Scientist, Meta) reposted the paper: “The evidence is accumulating.” [link]
- **Kevin Murphy** (Research Scientist, Google DeepMind; author of *Machine Learning: A Probabilistic Perspective*) reposted with his analysis of compounding errors in long reasoning chains. [link]
- **Vitaly Kurin** (Senior Research Scientist, NVIDIA): “Best paper I’ve read in a while. No excessive hype, clear question, rigorous empirical evaluation. Highly recommended.” [link]
- **YouTube** — in-depth analyses by *Automata Learning Lab* (331,131 views, [link]) and *mardin mardin* (55,797 views, [link]).

ProRL (NeurIPS 2025)

- **Twitter/X** — post about the paper reached **177,000 views**. [link]
- **Nathan Lambert** (reinforcement learning researcher, Allen Institute for AI; author of *The RLHF Book*) endorsed ProRL, highlighting its RL scaling results on a 1.5B model. [link]

SERVICE, TEACHING & OTHER AWARDS

Professional Service

- **Area Chair** — NAACL 2024; ACL Rolling Review (ARR), July 2025 and October 2023 (EACL 2024 / NAACL 2024 / ACL 2024 review cycle)
- **Reviewer** — ICML 2026, ICLR 2026, NeurIPS 2025, NeurIPS 2022, EMNLP 2022, and the Workshop on Commonsense Representation and Reasoning (CSRR) at ACL 2022
- **Journal Reviewer** — Journal of Artificial Intelligence Research (JAIR)
- **Panelist** — Machine Unlearning Challenge at NeurIPS 2023 [link]

Teaching Experience

- **Winter 2024** — Teaching Assistant, CSE 447/517 (Undergraduate/Graduate NLP), University of Washington
- **Winter 2021** — Teaching Assistant, CSE P517 (Professional NLP), University of Washington

Other Awards & Scholarships

- 2023** Best Senior Thesis Award, Paul G. Allen School of CSE [link]
- 2020** Outstanding Undergraduate Researcher Award, Runner-Up, Computing Research Association [link]
Lisa Simonyi Prize, Paul G. Allen School of CSE [link]
Levinson Emerging Scholars Award, University of Washington [link]
Mary Gates Research Scholarship, University of Washington
- 2019** Denton, Denice Dee Scholars Endowment, Paul G. Allen School of CSE
- 2018** Second Prize, UW Datathon (Citadel Investment Group)
Conference Travel Award, University of Washington